

# 7.3 Concurrency



This section has been set as optional by your instructor.

## Locking

The **concurrency system** is a database component that manages concurrent transactions.

The concurrency system implements isolation levels while attempting to optimize overall database performance. Tradeoffs between isolation levels and performance are complex, and many techniques have been implemented in concurrency systems. The most common technique is locking.

A **lock** is permission for one transaction to read or write data. Concurrent transactions are prevented from reading or writing the same data. A transaction takes locks when the transaction needs to read or write data. A transaction releases locks when the transaction is committed or no longer needs the locked data.

A **shared lock** allows a transaction to read, but not write, data. An **exclusive lock** allows a transaction to read and write data. Concurrent transactions can hold shared locks on the same data. When one transaction holds an exclusive lock, however, no concurrent transaction can take a shared or exclusive lock on the same data.

**Lock scope** is the collection of data reserved by a lock. Lock scope is often a single row, allowing other transactions to access other rows in the same table. If a transaction needs access to multiple rows, lock scope might be a block or the entire table. Since transactions also read and write indexes, lock scope might be an index entry, index block, or entire index.

The concurrency system monitors active transactions, determines when locks are needed, and issues requests to grant and release locks. Requests are sent to the **lock manager**, a component of the concurrency system that tracks, grants, and releases locks.

By requesting shared and exclusive locks as needed, the concurrency system implements the isolation level specified for each transaction. By minimizing lock scope and duration, the concurrency system reduces the duration of transactions that are waiting for locked data.

PARTICIPATION  
ACTIVITY

7.3.1: Shared and exclusive locking.



1 2 3 ←  2x speed

row locks

T <sub>1</sub>	T <sub>2</sub>	626	United Airlines	6:00
----------------	----------------	-----	-----------------	------

```
SELECT AirlineName
FROM Flight
WHERE FlightNumber = 702;
```

```
SELECT AirportCode
FROM Flight
WHERE FlightNumber = 702;
```

140	Aer Lingus	10:30	D
799	Air China	3:20	L
698	Air India	3:15	M
552	Lufthansa	18:00	N
107	Air China	3:20	A
971	Air India	19:35	S
702	American Airlines	5:05	M
1154	American Airlines	13:50	C
803	British Airways	13:50	A
1222	Delta Airlines	8:25	L
1001	Southwest Airlines	5:05	S

**T<sub>1</sub>****T<sub>2</sub>**

```
UPDATE Flight
SET AirlineName = 'Lufthansa'
WHERE FlightNumber = 107;
```

```
SELECT AirportCode
FROM Flight
WHERE FlightNumber = 107;
```

block locks

**T<sub>1</sub>****T<sub>2</sub>**

```
UPDATE Flight
SET AirlineName = 'Lufthansa'
WHERE FlightNumber = 803;
```

```
SELECT AirportCode
FROM Flight
WHERE FlightNumber = 1222
```

shared locks

exclusive locks

If T<sub>1</sub> takes an exclusive lock on a block, T<sub>2</sub> cannot access any row in the block until the lock is released.

Captions 

1. The concurrent transactions T<sub>1</sub> and T<sub>2</sub> can take a shared lock on the same row. The shared lock allows the transactions to read the row with FlightNumber 702.
2. If T<sub>1</sub> takes an exclusive lock on the row with FlightNumber 107, T<sub>2</sub> cannot access the row until the lock is released.
3. If T<sub>1</sub> takes an exclusive lock on a block, T<sub>2</sub> cannot access any row in the block until the lock is released.

**Feedback?**
**PARTICIPATION ACTIVITY**

7.3.2: Locking.



- 1) Several transactions can hold concurrent shared locks on the same row.

 True

 False
**Correct**

Shared locks are taken when transactions read data. Any number of transactions can read the same data concurrently.



- 2) Several transactions

**Correct**

- can hold concurrent exclusive locks on the same row.
- True  
 False

Exclusive locks are taken when transactions write data. Conflicts may arise when multiple transactions write the same data concurrently, so the lock manager grants an exclusive lock on a given row to one transaction at a time.

- 3) One transaction can hold an exclusive lock while other transactions hold shared locks on the same row.
- True  
 False

**Correct**



Conflicts may arise when one transaction reads data while another is writing the data, so the lock manager does not grant shared locks on a given row while any transaction holds an exclusive lock.

- 4) Several transactions can hold concurrent exclusive locks on the same block, as long as the transactions access different rows in the block.

- True  
 False

**Correct**



A block-level lock restricts access to all rows on the block. While a transaction holds an exclusive lock on a block, no other transaction can access any row on the block.

- 5) When a transaction takes a shared lock on a block, other transactions may be delayed.
- True  
 False

**Correct**



While a transaction holds a shared lock on a block, other transactions may take shared locks but not exclusive locks. Transactions requesting exclusive locks on the block, or any row in the block, must wait for release of the shared lock.

[Feedback?](#)

## Two-phase locking

When a transaction's isolation level is set to SERIALIZABLE, a transaction is fully isolated from other transactions. **Two-phase locking** is a specific locking technique that ensures serializable transactions. Three variations of two-phase locking are common in relational databases:

- **Basic two-phase locking** has expand and contract phases for each transaction, also known as grow and shrink phases. In the expand phase, the transaction can take, but not release, locks. In the contract phase, the transaction can release, but not take, locks.
- **Strict two-phase locking** holds all exclusive locks until the transaction commits or rolls back. The expand phase is the same as in basic two-phase locking, but the contract phase releases only shared locks.
- **Rigorous two-phase locking** holds both shared and exclusive locks until the transaction commits or rolls back. In effect, rigorous two-phase locking has no contract phase.

All three variations prevent conflicts and ensure serializable transactions. Strict and rigorous also prevent cascading rollbacks while basic does not. Rigorous is easier to implement than strict, but less efficient, since rigorous holds shared locks longer.

Concurrency systems in most relational databases implement strict two-phase locking.

**PARTICIPATION ACTIVITY**

## 7.3.3: Strict two-phase locking.



1 2 3 4 ←  2x speed

**Cascading Schedule**

<b>T<sub>1</sub></b>	<b>T<sub>2</sub></b>
write X	read X
rollback	<i>must rollback</i>

**Strict Two-Phase Locking**

<b>T<sub>1</sub></b>	<b>T<sub>2</sub></b>
write X	read X
rollback	<i>continues</i>

hold exclusive lock

wait for lock release

Rollback of T<sub>1</sub> does not cascade to T<sub>2</sub>.

Captions

1. Rollback of T<sub>1</sub> forces rollback of T<sub>2</sub>, so the schedule is cascading.
2. With strict two-phase locking, exclusive locks are not released during the contract phase. The exclusive lock is held until T<sub>1</sub> rolls back.

3. T<sub>2</sub> read must wait until the exclusive lock is released.
4. Rollback of T<sub>1</sub> does not cascade to T<sub>2</sub>.

[Feedback?](#)**PARTICIPATION ACTIVITY****7.3.4: Two-phase locking.**

- 1) Which two-phase locking technique results in the longest wait times for concurrent transactions?

- Basic
- Strict
- Rigorous

**Correct**

Rigorous holds both shared and exclusive locks until a transaction commits or rolls back. Any other transaction requesting locks on the same data must wait until commit or rollback.



- 2) Which two-phase locking technique has, in effect, just one phase?

- Basic
- Strict
- Rigorous

**Correct**

Rigorous holds all locks until a transaction commits or rolls back. Although rigorous is considered a type of two-phase locking, rigorous has only an expand phase.



- 3) A database uses basic two-phase locking. Transaction A takes an exclusive lock on X. Transaction B requests a shared lock on X. When is the shared lock granted?

- During the contract phase of A
- Immediately, when B requests the shared lock
- After A commits or rolls back

**Correct**

With basic two-phase locking, the exclusive lock is released during the contract phase of A. B is granted a shared lock when the exclusive lock is released.



OPTIONS BACK

- 4) A database uses strict two-phase locking. Transaction A takes an exclusive lock on X. Transaction B requests a shared lock on X. When is the shared lock granted?

- During the contract phase of A
- Immediately, when B requests the shared lock
- After A commits or rolls back

**Correct**

With strict two-phase locking, all exclusive locks are released at the time of commit or rollback. B's shared lock request must wait until A commits or rolls back.

**Feedback?**

## Deadlock

**Deadlock** is a state in which a group of transactions are frozen. In a deadlock, all transactions request access to data that is locked by another transaction in the group. None of the transactions can proceed. Ex:

1.  $T_1$  takes an exclusive lock on X.
2.  $T_2$  takes an exclusive lock on Y.
3.  $T_1$  requests a shared lock on Y.
4.  $T_2$  requests a shared lock on X.

$T_1$  waits for  $T_2$  to release the lock on Y.  $T_2$  waits for  $T_1$  to release the lock on X. The transactions are deadlocked.

A **dependent transaction** is waiting for data locked by another transaction. A **cycle** of dependent transactions indicates deadlock has occurred. Ex:  $T_1$  depends on  $T_2$ , which depends on  $T_3$ , which depends on  $T_1$ . The transactions are deadlocked.

Concurrency systems manage deadlock with a variety of techniques:

- **Aggressive locking.** Each transaction requests all locks when the transaction starts. If all locks are granted, the transaction runs to completion. If not, the transaction waits until other transactions release locks. Either way, the transaction cannot participate in a deadlock.

- **Data ordering.** All data needed by concurrent transactions is ordered, and each transaction takes locks in order. Taking locks in order prevents deadlock. Ex: The sequence above is modified so that locks on X precede locks on Y:

1. T<sub>1</sub> takes an exclusive lock on X.
2. T<sub>2</sub> requests a shared lock on X.
3. T<sub>1</sub> takes a shared lock on Y.
4. T<sub>2</sub> takes an exclusive lock on Y.

Request 2 waits for T<sub>1</sub> to release the exclusive lock on X. T<sub>1</sub> proceeds and eventually releases the lock on X. Now request 2 is granted and T<sub>2</sub> proceeds.

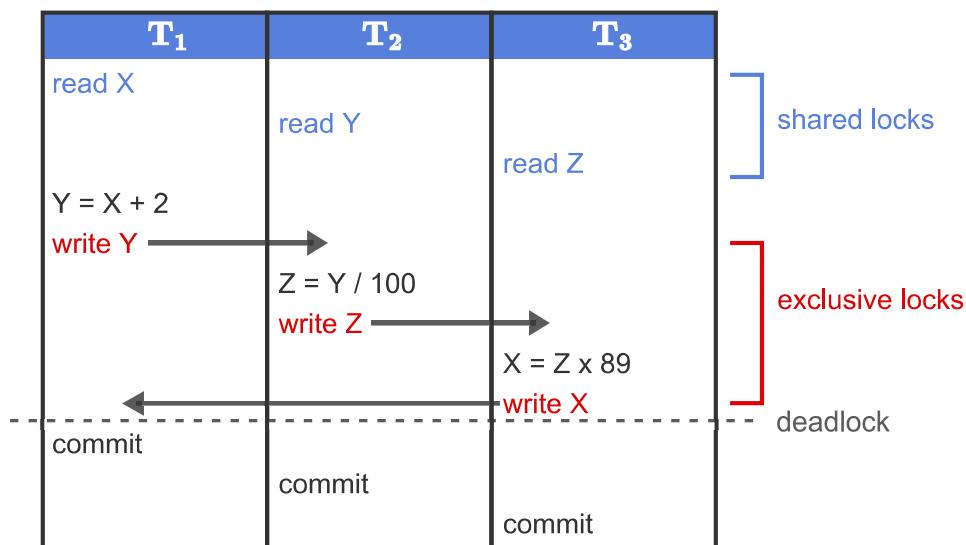
- **Timeout.** When waiting time for a lock exceeds a fixed period, the transaction requesting the lock rolls back. Alternatively, the concurrency system compares transaction start times and rolls back the later transaction. The timeout period is set by the database or configured by the database administrator.
- **Cycle detection.** The concurrency system periodically checks for cycles of dependent transactions. When a cycle is detected, the concurrency system selects and rolls back the 'cheapest' transaction. The cheapest transaction might, for example, have the fewest rows locked or most recent start time. The rollback breaks the deadlock.

PARTICIPATION  
ACTIVITY

7.3.5: Cycle of dependent transactions.



- 1 2 3 4 5 ← ✓ 2x speed



The dependency cycle causes deadlock.

Captions ▲

1. The schedule has three transactions that read and write X, Y, and Z.

2. Transactions first take shared locks when reading, then request exclusive locks when writing.
3. Since  $T_2$  has a shared lock on Y, the  $T_1$  exclusive lock request for Y waits for  $T_2$  to commit and release the shared lock.
4. Similarly,  $T_2$  waits for  $T_3$  to release the shared lock on Z, and  $T_3$  waits for  $T_1$  to release the shared lock on X.
5. The dependency cycle causes deadlock.

[Feedback?](#)
**PARTICIPATION ACTIVITY**

## 7.3.6: Deadlock.



- 1) Refer to the above animation. How many transactions must roll back to break the deadlock?

- One
- Two
- Three

**Correct**

Rolling back any one of the three transactions eliminates the cycle and breaks the deadlock. Ex: If  $T_3$  rolls back,  $T_2$  is no longer dependent and can proceed. When  $T_2$  commits,  $T_1$  can proceed.



- 2) Whenever deadlock occurs, a cycle of dependent transactions always exists.

- True
- False

**Correct**

Deadlock always involves a cycle. Without a cycle, some transaction is not dependent and can proceed, so deadlock cannot exist.



- 3) Deadlock can occur in a serializable schedule.

- True
- False

**Correct**

Serial schedules have no concurrent transactions and therefore cannot deadlock. Since a serializable schedule is equivalent to a serial schedule, serializable schedules also cannot deadlock.



- 4) A transaction with isolation level SERIALIZABLE can participate in a deadlock.

- True

**Correct**

In principle, SERIALIZABLE transactions run in serializable schedules and cannot deadlock. In practice, however, databases often enforce SERIALIZABLE transactions by allowing deadlock to occur and then rolling back one of the participating transactions.



False

- 5) Which deadlock management technique delays the fewest possible transactions?

 Aggressive locking Data ordering Timeout Cycle detection**Correct**

Cycle detection does not unnecessarily delay any transactions. Under cycle detection, transactions roll back only when participating in deadlock.

[Feedback?](#)

## Snapshot isolation

Exclusive locks prevent concurrent transactions from accessing data and therefore increase transaction duration. For this reason, many databases support alternative concurrency techniques. **Optimistic techniques** execute concurrent transactions without locks and, instead, detect and resolve conflicts when transactions commit. Optimistic techniques are effective when conflicts are infrequent, as in analytic applications with many reads and few updates.

One leading optimistic technique is snapshot isolation. **Snapshot isolation** creates a private copy of all data accessed by a transaction, called a **snapshot**, as follows:

1. Create a snapshot when the transaction starts.
2. Apply updates to the snapshot rather than the database.
3. Prior to commit, check for conflicts. A conflict occurs when concurrent transactions write the same data. Reads do not cause conflicts.
4. If no conflict is detected, write snapshot updates to the database (commit).
5. If a conflict is detected, discard the snapshot and restart the transaction (roll back).

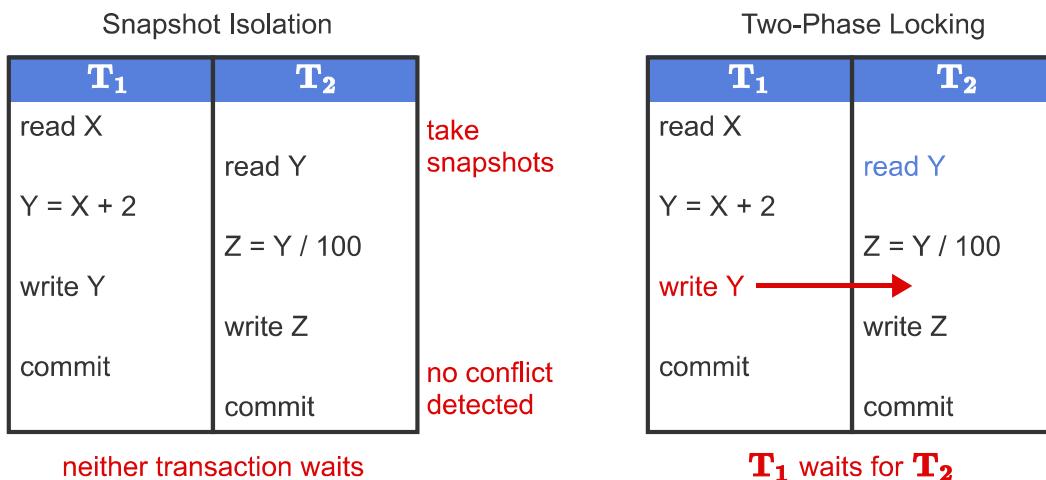
Since snapshot isolation does not take locks, transactions never wait. However, transactions occasionally restart after all operations are processed.

Two-phase locking is the most common concurrency technique, implemented in most relational databases. Snapshot isolation is a widely used alternative, available in Oracle and SQL Server.

**PARTICIPATION ACTIVITY**

7.3.7: Snapshot isolation vs. two-phase locking.





Under two-phase locking, **T<sub>1</sub>** waits for **T<sub>2</sub>** to commit and release the shared lock on Y.

Captions ^

1. Under snapshot isolation, the transactions take snapshots and process in parallel.
2. Since the transactions write different data, no conflict is detected. Both transactions commit.
3. Under snapshot isolation, neither transaction waits.
4. Under two-phase locking, **T<sub>1</sub>** waits for **T<sub>2</sub>** to commit and release the shared lock on Y.

[Feedback?](#)

When isolation level is set to SERIALIZABLE, two-phase locking ensures serializable schedules. Snapshot isolation does not ensure serializable schedules and may generate unexpected results. As a result, some databases implement a modified version of snapshot isolation.

**Serializable snapshot isolation** extends standard snapshot isolation and ensures serializable schedules when isolation level is set to SERIALIZABLE. Postgres supports serializable snapshot isolation.

PARTICIPATION  
ACTIVITY

7.3.8: Snapshot isolation allows non-serializable schedule.

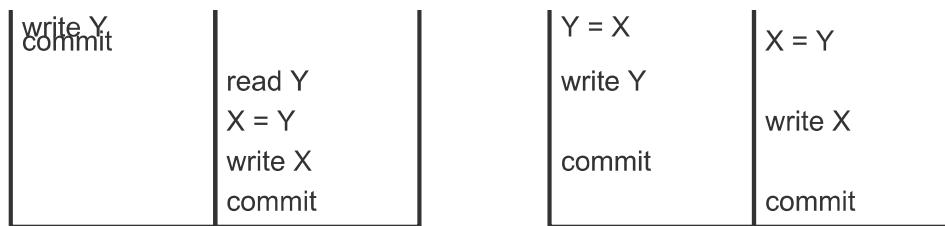


Serial Schedule

T <sub>1</sub>	T <sub>2</sub>
read X $Y = X$	

Non-Serializable Schedule

T <sub>1</sub>	T <sub>2</sub>
read X	read Y



**transactions commit  
X, Y have same value**

**Two-Phase Locking  
deadlock**

**Snapshot Isolation  
transactions commit  
X, Y swap values**

Snapshot isolation results in different values for X and Y.

Captions ^

1. The serial schedule commits and results in the same value for X and Y.
2. The non-Serializable schedule contains conflicting operations in a different order than the Serializable schedule.
3. The non-Serializable schedule contains a cycle and results in deadlock under two-phase locking.
4. Under snapshot isolation, no conflict is detected, and the transactions commit.
5. Snapshot isolation results in different values for X and Y.

[Feedback?](#)

PARTICIPATION  
ACTIVITY

7.3.9: Snapshot isolation.



Place snapshot isolation steps in the correct order.

**Make a private copy of data accessed by the transaction**

**Step 1**

Snapshot isolation begins by making a private copy of all data read and written by the transaction.

**Correct**

**Write updates to a private copy of data**

**Step 2**

During transaction execution, updates are applied to a private copy of data rather than the database, allowing the transaction

**Correct**

to proceed without waiting for locks.

### Determine if any updates conflict with other transactions

**Correct**

### Step 3

Prior to saving updates to the database, snapshot isolation checks if any other transactions updated the same data after the private copy was made.

### Write updates to the database or roll back the transaction

**Correct**

### Step 4

If updates do not conflict with other transactions, updates in the private copy are saved in the database. If updates conflict, the transaction rolls back.

**Reset**

**Feedback?**

**CHALLENGE ACTIVITY**

7.3.1: Concurrency.



379958.2369558.qx3zqy7

[Jump to level 1](#)



1



2



3



4

Consider the following concurrent transactions:

<b>T<sub>1</sub></b>	<b>T<sub>2</sub></b>
read Z	
X = Z * 41	read Y
write X	X = Y - 11
commit	write X
	commit

What occurs under:

snapshot isolation?

(c) T1 writes X, commits, and T2 rolls back.



rigorous two-phase locking?

- (a) T1 writes X, commits, and T2 waits until T1 commits. ▾

1

2

3

4

[Check](#)[Next](#)

**Done.** Click any level to practice more. Completion is preserved.



Expected:

- (c) T1 writes X, commits, and T2 rolls back.  
(a) T1 writes X, commits, and T2 waits until T1 commits.

Under snapshot isolation:

Transactions create a private copy of data when starting and check for write conflict. Since **T<sub>1</sub>** writes X first, **T<sub>1</sub>** writes X and commits. **T<sub>2</sub>** tries to write X second, so a conflict occurs and **T<sub>2</sub>** rolls back.

Under rigorous two-phase locking:

A transaction holds all shared and exclusive locks until the transaction commits or rolls back. **T<sub>1</sub>** starts with a shared lock on X before **T<sub>2</sub>** starts. **T<sub>1</sub>** immediately takes an exclusive lock on X. **T<sub>2</sub>** must wait until **T<sub>1</sub>** commits to take an exclusive lock on X.

[Feedback?](#)

How

was this

[Provide feedback](#)

section?